

Note and Formula of DDA4210

softstar

February 7, 2026

1 Lec1: introduction

There are no formulas in Lecture 1. Good luck with the rest!

2 Lec2: Advanced Ensemble Learning

2.1 Gradient Boosting:

Boosting is an ensemble technique that combines multiple weak learners to create a strong learner. The main idea is to train models sequentially, with each model focusing on the errors made by the previous ones. (用人话说，就是一波接一波地训练模型，每一波都专注于纠正前一波的错误，从而逐步提升整体的预测能力。)

- **Final Model:**

$$H_T(\mathbf{x}) = \sum_{t=1}^T \alpha_t h_t(\mathbf{x})$$

- The weighted sum of all weak learners(加权和).
- $H_T(\mathbf{x})$ is the final model after T rounds.
- $h_t(\mathbf{x})$ is the t -th weak learner.
- α_t is the weight of the t -th weak learner.

- **Loss Function:**

$$\mathcal{L}(H) := \frac{1}{n} \sum_{i=1}^n \ell(H(\mathbf{x}_i), y_i)$$

- **Each Step:** We want to add a function h to minimize the loss as fast as possible. Using first-order Taylor expansion(一阶泰勒展开):

$$\mathcal{L}(H + \alpha h) \approx \mathcal{L}(H) + \alpha \langle \nabla \mathcal{L}(H), h \rangle$$

- **Find h :** Minimize $\langle \nabla \mathcal{L}(H), h \rangle$, i.e.:

$$h = \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n \frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]} h(\mathbf{x}_i)$$

* Here $\frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]}$ is the gradient of the loss function w.r.t. the current model's output on the i -th sample.

* We train h to fit these **negative gradients**(负梯度).

- **GBR with squared error loss:** For regression tasks with squared error loss:

$$\ell(H(\mathbf{x}_i), y_i) = \sum_{i=1}^n (y_i - H(\mathbf{x}_i))^2$$

- Solve $h_{t+1} = \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n q_i h(\mathbf{x}_i)$, where $q_i = \frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]}$
- Let $\sum_{i=1}^n h^2(\mathbf{x}_i) = \text{constant}$ (we can always normalize the predictions(对结果归一化处理)) and replace q_i with $-2r_i$. We have

$$\begin{aligned} h_{t+1} &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n q_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} -2 \sum_{i=1}^n r_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n (r_i^2 - 2r_i h(\mathbf{x}_i) + (h(\mathbf{x}_i))^2) \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n (h(\mathbf{x}_i) - r_i)^2 \end{aligned} \tag{1}$$

- We train h_{t+1} to predict r_i , which are from the old model H_t .

- The gradient is:

$$\frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]} = -2(y_i - H(\mathbf{x}_i))$$

- So we fit h to the residuals:

$$r_i = y_i - H(\mathbf{x}_i)$$

- Update the model:

$$H_t(\mathbf{x}) = H_{t-1}(\mathbf{x}) + \alpha_t h_t(\mathbf{x})$$

- **GBR with Absolute loss (更鲁棒):** For regression tasks with absolute loss:

- Square loss is easy to deal with mathematically but not robust to outliers.
- Absolute loss (more robust to outliers):

$$\ell(y, \hat{y}) = |y - \hat{y}|$$

- The gradient is $q_i = \frac{\partial \mathcal{L}}{\partial H(\mathbf{x}_i)} = -\text{sign}(y_i - H(\mathbf{x}_i))$.
- Then fit h on $-q_i$, $i = 1, 2, \dots, n$. (no longer the residuals, different from using the squared loss)

- **GBR with Huber loss(更更鲁棒):**

- Huber loss (more robust to outliers):

$$\ell(y, \hat{y}) = \begin{cases} \frac{1}{2}(y - \hat{y})^2 & |y - \hat{y}| \leq \delta \\ \delta(|y - \hat{y}| - \delta/2) & |y - \hat{y}| > \delta \end{cases}$$

- The gradient is:

$$\frac{\partial \mathcal{L}}{\partial H(\mathbf{x}_i)} = \begin{cases} -(y_i - H(\mathbf{x}_i)) & |y_i - H(\mathbf{x}_i)| \leq \delta \\ -\delta \operatorname{sign}(y_i - H(\mathbf{x}_i)) & |y_i - H(\mathbf{x}_i)| > \delta \end{cases}$$

- **GBM for classification:**

- Predict probability of K classes:

$$p_k(\mathbf{x}) = \frac{\exp(h^{(k)}(\mathbf{x}))}{\sum_{c=1}^K \exp(h^{(c)}(\mathbf{x}))} \triangleq \hat{y}^{(k)}, \quad k = 1, 2, \dots, K$$

- Loss: $\mathcal{L}(H) = \sum_{i=1}^n \ell(\mathbf{y}_i, \hat{\mathbf{y}}_i)$ (e.g. cross-entropy or KL divergence)

- Initialize $H^{(1)}, H^{(2)}, \dots, H^{(K)}$, iterate until converge or reach max T :

1. Calculate negative gradients for every class:

$$-g_k(\mathbf{x}_i) = -\frac{\partial \mathcal{L}}{\partial [H^{(k)}(\mathbf{x}_i)]}, \quad i = 1, \dots, n, \quad k = 1, \dots, K$$

2. Fit $h^{(k)}$ to $-g_k(\mathbf{x}_i)$ (负梯度), $k = 1, 2, \dots, K$.

3. Update: $H^{(k)} \leftarrow H^{(k)} + \alpha h^{(k)}$, $k = 1, 2, \dots, K$.

2.2 AdaBoost

A special case of gradient boosting with exponential loss.

- Exponential loss (learns α adaptively):

$$\mathcal{L}(H) = \sum_{i=1}^n e^{-y_i H(\mathbf{x}_i)}$$

- Gradient: $q_i = \frac{\partial \mathcal{L}}{\partial H(\mathbf{x}_i)} = -y_i e^{-y_i H(\mathbf{x}_i)}$
- Let $w_i = \frac{1}{Z} e^{-y_i H(\mathbf{x}_i)}$, where $Z = \sum_{i=1}^n e^{-y_i H(\mathbf{x}_i)}$ (constant w.r.t h), so $\sum_{i=1}^n w_i = 1$. w_i is the relative contribution of (\mathbf{x}_i, y_i) to the overall loss.
- Binary classification: $y \in \{-1, +1\}$, $h(\mathbf{x}) \in \{-1, +1\}$.
- Derivation:

$$\begin{aligned} h &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n q_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} - \sum_{i=1}^n w_i y_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i - \sum_{i:h(\mathbf{x}_i) = y_i} w_i \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i \end{aligned}$$

* Last equality holds because $\sum_{i=1}^n w_i = 1$.

- Result:

$$h = \arg \min_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i$$

$$\epsilon := \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i$$

is the weighted classification error(加权错误率).

Note: misclassified points by H get larger weights(分类错误的点会得到更大的权重).

- Given h , find α via:

$$\alpha = \arg \min_{\alpha} \mathcal{L}(H + \alpha h) = \arg \min_{\alpha} \sum_{i=1}^n e^{-y_i(H(\mathbf{x}_i) + \alpha h(\mathbf{x}_i))}$$

- Differentiate w.r.t α and set to zero:

$$\sum_{i=1}^n y_i h(\mathbf{x}_i) e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} = 0$$

It follows that:

$$\begin{aligned} \sum_{i:h(\mathbf{x}_i)y_i=1} e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} - \sum_{i:h(\mathbf{x}_i)y_i=-1} e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} &= 0 \\ \sum_{i:h(\mathbf{x}_i)y_i=1} w_i e^{-\alpha} - \sum_{i:h(\mathbf{x}_i)y_i=-1} w_i e^{\alpha} &= 0 \end{aligned}$$

We have $(1 - \epsilon)e^{-\alpha} - \epsilon e^{\alpha} = 0$, $e^{2\alpha} = \frac{1-\epsilon}{\epsilon}$, and get:

$$\boxed{\alpha = \frac{1}{2} \ln \frac{1-\epsilon}{\epsilon}}$$

2.3 Mixture of Experts(MoE)

A machine learning technique where multiple expert learners (e.g. neural networks) are used to divide a problem space into homogeneous regions (distinct subtasks). (用人话 (AI) 说就是，把一个复杂的问题拆分成多个子任务，每个子任务由一个专家模型来处理，从而提升整体的学习效果。)

(骗你的人话也没看懂。。)

2.3.1 The First Attempt

- Error function:

$$E = \left\| y - \sum_{j=1}^k g_j O_j \right\|^2$$

- y : target vector; O_j : output of expert j ; g_j : proportional contribution of expert j .

- *This error function does not ensure localisation of experts.(人话：专家没有明确的分工).

2.3.2 The Second Attempt

- Error function:

$$E = \sum_{j=1}^k g_j \|y - O_j\|^2$$

- The system tends to devote a single expert to each training case.
- *This may not work well in practice.(人话: 实际效果可能不佳)

2.3.3 The Third Attempt [Jacobs et al. 1991]

- Error function (mixture of Gaussians):

$$E_{ME} = -\log \sum_{j=1}^k g_j \exp \left(-\frac{1}{2}(y - O_j)^T \Sigma^{-1}(y - O_j) \right)$$

- Assume $\Sigma = I$ (Σ : Covariance matrix, 协方差矩阵), derivative w.r.t the i -th expert:

$$\frac{\partial E_{ME}}{\partial O_i} = - \left[\frac{g_i \exp(-\frac{1}{2}(y - O_i)^T(y - O_i))}{\sum_j g_j \exp(-\frac{1}{2}(y - O_j)^T(y - O_j))} \right] (y - O_i)$$

- Compare with derivative of second attempt:

$$\frac{\partial E}{\partial O_i} = -2g_i(y - O_i)$$

- The former considers how well expert i performs relative to others, adapting the best-fitting expert faster.

E.g., $g_1 = 0.8$, $g_2 = 0.2$, then $\frac{0.8 \times 0.9}{0.8 \times 0.9 + 0.2 \times 0.1} \approx 0.97 > 0.8$.
(人话: 表现好的专家会得到更快的提升)

(这 nm 都是什么玩意?)

2.4 Stacking(堆叠法)

Multiple base learners' outputs $(\hat{y}_1, \hat{y}_2, \dots, \hat{y}_N)$ are combined by a meta-learner (stacker) to produce the final prediction \hat{Y} . (用人话说就是, 把多个模型的预测结果再拿去训练一个新的模型, 从而提升整体的预测效果。)

- Multi-level stacking: stacking can be repeated.
- Popular stackers: linear models (fast), gradient boosting (accurate).
- Base models should be diverse, expert at different data parts.
- Trade-off: accurate but slow to predict.

(为什么一讲能有这么多神秘小知识点和神秘公式?)

3 Lec3: Advanced Applications

3.1 Recommendation Systems

3.1.1 Collaborative Filtering(协同过滤)

Core idea: Use behavioral data from many users (e.g., ratings, clicks) to predict what the current user might like. (人话: 大数据推荐你喜欢的东西)

- **User-Item Interaction:**
 - Explicit Feedback(显式反馈): ratings, purchases.
 - Implicit Feedback(隐式反馈): clicks, browsing time.
- **User-Item Rating Matrix:** Rows = users, columns = items, values = ratings. Typically large and sparse.
- **Distance/Similarity Measurement(相似度度量):**
 - Euclidean distance: $\text{sim}(\text{user}_i, \text{user}_j) = \frac{1}{1 + \|\mathbf{x}_i - \mathbf{x}_j\|_2}$
 - Cosine similarity: $\text{sim}(\text{user}_i, \text{user}_j) = \frac{\mathbf{x}_i \cdot \mathbf{x}_j}{\|\mathbf{x}_i\| \|\mathbf{x}_j\|}$
 - Pearson correlation: $\rho_{X,Y} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y}$
- **Nearest-Neighbor Collaborative Filtering(最近邻协同过滤):**
Predict utility of item i based on similar users who rated that item.
 - \mathcal{N} : neighborhood set (most similar users to u who rated item i) (人话: 给你推荐东西的那些“邻居”用户)
 - $w_{uv} \in [0, 1]$: similarity weight between users u and v (人话: 你和邻居用户的相似度)
 - Prediction:

$$\hat{x}_{ui} = \bar{x}_u + \sum_{v \in \mathcal{N}} \left((x_{vi} - \bar{x}_v) \times \frac{w_{uv}}{\sum_{v' \in \mathcal{N}} w_{uv'}} \right)$$
 - $\bar{x}_u = \frac{1}{|I_u|} \sum_{i \in I_u} x_{ui}$ (average rating of user u , where I_u is the set of items rated by u)
 - $\bar{x}_v = \frac{1}{|I_v|} \sum_{i \in I_v} x_{vi}$ (average rating of user v , where I_v is the set of items rated by v)
- **Matrix Factorization Collaborative Filtering(矩阵分解协同过滤):**

- Notations:

- * $R = [r_{ui}] \in \mathbb{R}^{m \times n}$: incomplete user-item rating matrix
- * Ω : set of observed entries (known ratings)
- * $P = [p_1, \dots, p_u, \dots, p_m] \in \mathbb{R}^{f \times m}$, $Q = [q_1, \dots, q_i, \dots, q_n] \in \mathbb{R}^{f \times n}$

- Basic SVD ($R \approx P^T Q$):

$$\min_{P,Q} \sum_{(u,i) \in \Omega} \{(r_{ui} - p_u^T q_i)^2 + \lambda(\|p_u\|^2 + \|q_i\|^2)\}$$

- SVD with bias ($b_{ui} = \mu + b_u + b_i$):

$$\min_{P,Q,B} \sum_{(u,i) \in \Omega} \{(r_{ui} - \mu - b_u - b_i - p_u^T q_i)^2 + \lambda(\|p_u\|^2 + \|q_i\|^2 + b_u^2 + b_i^2)\}$$

- μ : global mean; b_u : user bias; b_i : item bias;
- $\lambda(\cdot)$: regularization term (prevents overfitting).
- Optimization: GD/SGD or Alternating Least Squares(交替最小二乘法).

- **Pros & Cons of Collaborative Filtering:**

- Pros: No domain knowledge needed; captures diverse user preferences(人话: 不需要领域知识, 能捕捉多样的用户偏好).
- Cons: Cold-start problem (new users/items); data sparsity(人话: 新来的用户和物品数据较少, 导致推荐效果差).

3.1.2 Content-Based Methods

- **Content analysis:** item \rightarrow feature vector v (e.g. TF-IDF, image features).
- **Profile learning:** user \rightarrow feature vector z (e.g. age, sex, education).
- **Filtering module:** train classification/regression model to predict user's utility for an item.
- **Recommendation for user:**
 - n : # of items;
 - $z_u \in \mathbb{R}^d$:
 - user u 's feature vector;
 - $h : \mathbb{R}^d \rightarrow \mathbb{R}^n$ (e.g. neural network);
 - h_i : i -th output of h .
 - ℓ : loss function (e.g. squared loss).
$$\min_h \sum_{(u,i) \in \Omega} \ell(r_{ui}, h_i(z_u))$$
- **Recommendation for item:**
 - m : # of users;
 - $v_i \in \mathbb{R}^{d'}$: item
 - i 's feature vector;
 - $g : \mathbb{R}^{d'} \rightarrow \mathbb{R}^m$ (e.g. neural network);
 - g_u : u -th output of g .
 - ℓ : loss function (e.g. squared loss).
$$\min_g \sum_{(u,i) \in \Omega} \ell(r_{ui}, g_u(v_i))$$
- **Pros & Cons of Content-Based Methods:**
 - Pros: User-independent; explainable; handles new items/users well.
 - Cons: Needs domain knowledge; narrow recommendations (similar items).

3.1.3 Hybrid Methods(看起来是不重要的知识点)

Most modern systems are hybrid recommenders.

- Combine separate recommenders (CF + CB): ensemble techniques (linear weighting, stacking, etc.)
- Add content-based aspects to CF: e.g. matrix factorization with side information.

3.1.4 Evaluation Metrics for RS

3.1.4.1 Prediction Metrics(评分预测指标)

- **Mean Absolute Error (MAE):**

$$\text{MAE} = \frac{1}{|\mathcal{T}|} \sum_{(u,i) \in \mathcal{T}} |r_{ui} - \hat{r}_{ui}|$$

- **Root Mean Squared Error (RMSE):**

$$\text{RMSE} = \sqrt{\frac{1}{|\mathcal{T}|} \sum_{(u,i) \in \mathcal{T}} (r_{ui} - \hat{r}_{ui})^2}$$

Here \mathcal{T} denotes the set of user–item pairs used for evaluation (e.g., the test set).(人话：测试集中所有用户和物品的组合)

3.1.4.2 Ranking-based Metrics(基于排序的指标)

- **Precision@K:** fraction of top- K recommended items that are relevant(人话：前 K 个推荐中有多大是相关的).

$$\text{Prec}(R)_k = \frac{|\{r \in R : r \leq k\}|}{k}$$

- **Recall@K:** fraction of relevant items covered in top- K (人话：前 K 个推荐覆盖了多少相关的物品).

$$\text{Recall}(R)_k = \frac{|\{r \in R : r \leq k\}|}{|R|}$$

- (Precision = TP/(TP+FP); Recall = TP/(TP+FN)).
- r = rank position of a recommended item;
- k = cut-off (top- k);
- R = set of relevant items for the user.

- **Average Precision (AP@N):** average of precision values at ranks of relevant items.(前 N 个推荐中相关物品的精确率)

$$\text{AP}@N = \frac{1}{m} \sum_{k=1}^N P(k) \cdot \text{rel}(k)$$

where $P(k)$ is precision@ k , m number of relevant items, and $\text{rel}(k)$ is indicator if item at rank k is relevant.

- **Mean Average Precision (MAP):** mean of AP over Q users:

$$\text{MAP} = \frac{1}{Q} \sum_{q=1}^Q \text{AP}(q)$$

- **Normalized Discounted Cumulative Gain (NDCG):** evaluates ranked relevance with position discounting.

$$NDCG_p = \frac{DCG_p}{IDCG_p}, \quad DCG_p = \sum_{i=1}^p \frac{2^{rel_i} - 1}{\log_2(i + 1)}$$

where rel_i is relevance of item at rank i , and $IDCG_p$ is the ideal DCG (sorted by relevance).

Range: $[0, 1]$. (NDCG 越接近 1 越好)

- **Example:** 5 recommended items with relevances $[3, 2, 1, 0, 2]$ (in rank order).

$$DCG_5 = \frac{2^3 - 1}{\log_2(1 + 1)} + \frac{2^2 - 1}{\log_2(2 + 1)} + \frac{2^1 - 1}{\log_2(3 + 1)} + \frac{2^0 - 1}{\log_2(4 + 1)} + \frac{2^2 - 1}{\log_2(5 + 1)} \approx 10.5538,$$

$$IDCG_5 = DCG_5(\text{sorted rel} = [3, 2, 2, 1, 0]) \approx 10.8235,$$

$$NDCG_5 = \frac{DCG_5}{IDCG_5} \approx 0.975.$$

(这一坨又是什么玩意？)

3.2 Learning to Rank(L2R) (排序学习)

Learning to Rank (L2R) trains models to order items by relevance, optimizing ranking-specific objectives such as pairwise or listwise losses. (你说得对，但是这里好像也没有什么公式啊(雾))

- L2R is a supervised learning problem for ranking.
- Training data consists of:
 - A set of queries $Q = \{q_1, \dots, q_m\}$
 - A set of documents D
 - For each query i , relevant documents $D_i = \{d_{i,1}, \dots, d_{i,n_i}\} \subseteq D$
 - Relevance scores $\mathbf{y}_i = (y_{i,1}, \dots, y_{i,n_i})$ for each $d_{i,j}$
- Goal: Given a new query q , output a sorted list of relevant documents.
- **Point-wise Modeling:** Predicts each (query, document) pair independently.

Pro: Simple, can use regression/classification.
Con: Ignores relative order between documents.
 (一句话：点对点建模简单但忽略了文档间的相对顺序。)
- **Pair-wise Modeling:** Predicts preference between document pairs for a query.

Pro: Models relative order.
Con: Cannot distinguish excellent-bad from fair-bad pairs.
 (一句话：成对建模能捕捉相对顺序，但无法区分优秀-差和一般-差的对比。)
- **List-wise Modeling:** Predicts for the whole ranked list of documents.

Pro: Considers position in ranking, aligns with ranking metrics.
Con: High training complexity.
 (一句话：列表建模考虑排名位置，但训练复杂度高。)
- **Evaluation for L2R:** Use benchmark datasets and ranking metrics (e.g., MAP, NDCG).

- Algorithms for L2R:

- Example: Ranking SVM (pairwise):

- * **Goal:** Learn a scoring function $h(x) = w^\top x$ such that for any pair with labels $y_i > y_j$ we have $h(x_i) > h(x_j)$. (人话: 让相关性更高的文档得分更高)

- * **Training pairs:** Construct pair set $\mathcal{P} = \{(i, j) : y_i > y_j\}$; $m = |\mathcal{P}|$ denotes number of pairs.

- * **Optimization (primal):**

$$\begin{aligned} \min_{w, \xi_{ij} \geq 0} \quad & \frac{1}{2} \|w\|^2 + \frac{C}{m} \sum_{(i,j) \in \mathcal{P}} \xi_{ij} \\ \text{s.t.} \quad & w^\top x_i \geq w^\top x_j + 1 - \xi_{ij}, \quad \forall (i, j) \in \mathcal{P}. \end{aligned}$$

- * **Notes:** ξ_{ij} are hinge-loss slacks; C controls margin vs. training error. The objective is equivalent to minimizing the average pairwise hinge loss.

- * **Prediction:** Score each document by $h(x) = w^\top x$ and sort descending to produce a ranking.

- * **Remarks:** Works well for pairwise preferences; training can be expensive due to $O(n^2)$ pairs, so sampling or stochastic methods are often used. Kernel SVMs and regularization extend naturally.

- * (人话总结: Ranking SVM 通过学习一个线性评分函数来排序文档, 优化目标是最大化正确排序对的间隔 + 最小化排序错误的惩罚。训练时需要处理大量文档对, 复杂度高。)

(叽里咕噜说什么在)

4 Lec4-1: Graph Cut and Spectral Clustering(谱聚类)

4.1 Graph Partition

A similarity graph $G=(V,E,W)$ represents data points as vertices V , with an edge in E when the pairwise similarity is positive and weights W storing those affinities. The affinity matrix records these pairwise similarities. Graph partitioning (clustering) aims to split the graph so that edges inside a group have large weights while edges across groups have small weights. (人话: 图划分就是把图分成若干部分, 使得每个部分内的节点之间联系紧密 (组内权重大), 而不同部分之间的联系较弱 (组间权重小)。)

Given data points, a similarity graph can be constructed using methods such as **k-nearest neighbor** or ϵ -neighborhood. The edge weights are often defined by a **Gaussian kernel**:

$$k(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

4.2 Minimum Cut

- Minimum cut: partition the graph into two sets A and B minimizing $\text{cut}(A, B) := \sum_{i \in A, j \in B} w_{ij}$.
- Solvable efficiently (e.g. via max-flow/min-cut algorithms, typical cost $O(|V||E|)$), but the minimum cut often yields unbalanced solutions (it may isolate vertices). (人话: 最小割可以高效求解, 但结果往往不平衡, 可能会把一些节点孤立出来。)

- Not satisfactory partition? Often isolates vertices

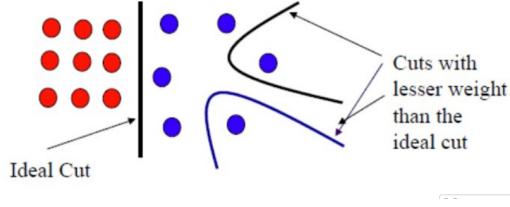


图 1: Minimum Cut Example

(画的不错的一张图)

- To address this, we can use **Normalized Cut (Ncut)**:

4.3 Normalized Cut

Normalized Cut balances cut weight with cluster sizes. For a partition (A, B) , define

$$\text{vol}(A) = \sum_{i \in A} d_i, \quad d_i = \sum_j w_{ij},$$

and

$$\text{Ncut}(A, B) := \text{cut}(A, B) \left(\frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right).$$

- Minimizing Ncut favors balanced partitions but is **NP-hard**;
- **spectral clustering (谱聚类)** provides an efficient relaxation.

4.3.1 Degree Matrix and Graph Laplacian

$$W = \begin{bmatrix} w_{11} & w_{12} & \cdots & w_{1N} \\ w_{21} & w_{22} & \cdots & w_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ w_{N1} & w_{N2} & \cdots & w_{NN} \end{bmatrix}$$

$$D = \begin{bmatrix} d_1 & 0 & \cdots & 0 \\ 0 & d_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & d_N \end{bmatrix}, \quad d_j = \sum_{i=1}^N w_{ij}.$$

- D is the **degree matrix**. And the **graph Laplacian matrix** is defined as

$$L := D - W.$$

- Properties: L is symmetric positive semi-definite, $\mathbf{1}^\top L = 0$, and its eigenvectors are used in spectral clustering (relaxation of Ncut).

4.3.2 Normalized Cut and Graph Laplacian

Mathematical derivation (optional)

Recall $L = D - W$ and $D = \text{diag}(d_1, \dots, d_N)$.

Let $u = [u_1, u_2, \dots, u_N]^\top$ with

$$u_i = \begin{cases} \frac{1}{\text{vol}(A)}, & \text{if } i \in A, \\ -\frac{1}{\text{vol}(B)}, & \text{if } i \in B. \end{cases}$$

Then

$$u^\top Lu = \frac{1}{2} \sum_{i,j} w_{ij} (u_i - u_j)^2 = \sum_{i \in A, j \in B} w_{ij} \left(\frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right)$$

and

$$u^\top Du = \sum_i d_i u_i^2 = \sum_{i \in A} \frac{d_i}{\text{vol}(A)^2} + \sum_{j \in B} \frac{d_j}{\text{vol}(B)^2} = \frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)}.$$

Therefore

$$\frac{u^\top Lu}{u^\top Du} = \sum_{i \in A, j \in B} w_{ij} \left(\frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right) = \text{Ncut}(A, B).$$

(这几把都是什么玩意？好像是上面的推导吧)

(不管了，反正是 optional，摆在这图个吉利)

Conclusion

- Ncut is equivalent to minimizing the Rayleigh quotient $\frac{u^\top Lu}{u^\top Du}$, i.e.,

$$\min_{A,B} \text{Ncut}(A, B) \iff \min_u \frac{u^\top Lu}{u^\top Du}, \quad u \in \mathbb{R}^N, \quad u_i = \begin{cases} \frac{1}{\text{vol}(A)}, & i \in A, \\ -\frac{1}{\text{vol}(B)}, & i \in B. \end{cases}$$

- Equivalent formulation: minimize the quotient subject to $u^\top D\mathbf{1} = 0$ and binary constraints $u_i \in \{1, -b\}$ (for some $b > 0$).
- Relaxation:

$$Lu = \lambda Du,$$

taking the eigenvector corresponding to the **second smallest eigenvalue** as the relaxed solution.

- Equivalently, use the normalized Laplacian $\tilde{L} = D^{-1}L = I - D^{-1}W$.
Obtain a binary partition by thresholding u at 0: $i \in A$ if $u_i \geq 0$, else $i \in B$.
- Extend to k clusters by using the first k nontrivial eigenvectors and applying k -means (spectral clustering).
- (人话：2类划分找第二小特征值对应的特征向量， k 类划分用前 k 个非平凡特征向量，然后 k -means 聚类。)

(事实上还是没看懂，插个眼以后复习的时候看看有没有什么需要补的)

4.4 Spectral Clustering Algorithm

- Input: data $X = \{x_1, x_2, \dots, x_N\}$ and number K of clusters.
- **Step 1:** Construct a similarity (affinity) matrix W (e.g. Gaussian kernel)

$$w_{ij} = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

and build either a k -nearest neighbor or ϵ -neighborhood graph.

- **Step 2:** Compute Laplacian L (or normalized variants):

$$L = D - W, \quad \tilde{L} = D^{-1}L = I - D^{-1}W, \quad \hat{L}_{sym} = I - D^{-1/2}WD^{-1/2}.$$

- **Step 3:** Eigen-decompose (normalized) Laplacian and take first K nontrivial eigenvectors to form $Z = [v_1, \dots, v_K]^\top \in \mathbb{R}^{K \times N}$.(对 L 做特征值分解)
- **Step 4:** Normalize the columns (row-wise embedding) to unit ℓ_2 norm(归一化):

$$z_i \leftarrow z_i / \|z_i\|, \quad i = 1, \dots, N.$$

- **Step 5:** Run K -means on $\{z_1, \dots, z_N\}$ and output K clusters (assignments on Z or map back to X).
- Notes: use \tilde{L}_{sym} (symmetric normalized) for best numerical stability; thresholding the second eigen-vector recovers a binary partition.

- **Properties of L:**

For $L = D - W$ or $\hat{L} = I - D^{-1/2}WD^{-1/2}$

- L (and \hat{L}) are symmetric and positive semi-definite.(对称 + 半正定)
- Eigenvalues satisfy $0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_N$.
- The multiplicity K of eigenvalue 0 equals the number of connected components of the graph (hence K clusters). (0 特征值的数量等于图的连通分量数, 也就是聚类数)

5 Lec4-2: Semi-Supervised Learning (SSL, 半监督学习)

5.1 Notations

- Input (or feature) $\mathbf{x} \in \mathcal{X}$, output (or label) $\mathbf{y} \in \mathcal{Y}$
- Learner $f : \mathcal{X} \rightarrow \mathcal{Y}$
- Labeled data $(X_l, Y_l) = \{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_l, \mathbf{y}_l)\}$
- Unlabeled data $X_u = \{\mathbf{x}_{l+1}, \dots, \mathbf{x}_N\}$, available during training
- Loss function $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$
- Usually, $l \ll N$
- Test data $X_{\text{test}} = \{\mathbf{x}_{N+1}, \dots\}$, not available during training

Why Semi-Supervised Learning?

Labeled data is often **scarce**, **expensive**, and requires expert annotation or specialized equipment. In contrast, unlabeled data is **abundant and cost-effective**. By leveraging both, we can build more robust and high-performing machine learning models.

5.2 Self-Training Algorithm

Assumption: One's own high confidence predictions are correct.(人话：自己高置信度的预测是正确的)

5.2.1 Algorithm Procedure:

1. Train model f from labeled data (X_l, Y_l) .
2. Predict labels for unlabeled data $x \in X_u$. (用有标签数据训练模型，然后对无标签数据进行预测)
3. Add the pseudo-labeled pairs $(x, f(x))$ to the labeled training set. (把预测结果当作新的标签加入训练集)
4. Repeat the process until a stopping criterion is met.

5.2.2 Variations of Self-Training

Depending on how pseudo-labeled data is incorporated, there are several variations:

- **Most Confident:** Add only a few samples $(x, f(x))$ with the highest prediction confidence to the labeled data. (人话：只添加那些模型最有信心的预测结果)
- **All Samples:** Add all predicted $(x, f(x))$ directly to the labeled data.(直接把所有预测结果都加入训练集)
- **Weighted Approach:** Add all $(x, f(x))$ to the labeled data, but assign different weights to each sample according to its prediction confidence. (根据预测置信度给每个样本分配不同的权重)

5.2.3 Self-Training Algorithm: Propagating 1-NN (传播 1-NN)

This is a specific instance of self-training using the 1-Nearest Neighbor (1-NN) approach:

1. Classify an unlabeled sample x using the 1-NN rule (assigning the label of its nearest labeled neighbor).(用 1-NN 方法对无标签样本进行分类，分配最近的有标签邻居的标签)
2. Add the newly labeled pair $(x, f(x))$ to the labeled dataset and repeat the process.

Advantages and Disadvantages

- **Pros:** It is the simplest semi-supervised approach, serves as a versatile wrapper method, and is widely utilized in Natural Language Processing (NLP).
- **Cons:** The method is sensitive to outliers, and early classification errors can lead to self-reinforcing mistakes. (对异常值敏感)

5.3 Graph based SSL methods

Assumption: A graph is given on the labeled and unlabeled data. Instances connected by heavy edge tend to have the same label(由重边连接的实例往往具有相同的标签)

人话就是边的权重反映了两个节点(数据点)之间的相似程度，边权越大，数据点越相似

5.3.1 Graph Construction

- **Nodes:** The set of nodes is the union of labeled and unlabeled data, $X_l \cup X_u$.
- **Edges:** Constructed via a k -Nearest Neighbor graph (unweighted) or a fully connected graph where weights decay with distance (全连接图, 权重随距离衰减):

$$w_{ij} = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

- w_{ij} 是高斯核的那个函数，反映了数据点之间的相似度

Regularized Classifier A standard classifier learns by minimizing a loss term combined with regularization (e.g., LASSO or Regularized Least Squares). We can extend this to semi-supervised learning by using unlabeled data for regularization.

Graph-Based Regularization The core principle is that if two instances x_i and x_j are similar (indicated by a large edge weight w_{ij}), their predicted labels $f(x_i)$ and $f(x_j)$ should also be similar. This leads to the following optimization problem:

其实就是学习一个学习器来使得损失函数 + 正则化项最小

$$\min_f \sum_{i=1}^l \ell(y_i, f(x_i)) + \lambda \sum_{i=1}^N \sum_{j=1}^N w_{ij} \|f(x_i) - f(x_j)\|^2$$

第一项：有标签数据的损失。第二项：图正则化项，强制相似点有相似预测。 λ : 正则化参数。

Algorithm Objectives The graph-based algorithm operates with two primary goals:

- **Label Constraint:** It enforces the correct labels on the existing labeled data.
- **Manifold Consistency:** It maximizes the consistency of unlabeled examples relative to the underlying graph topology.

5.3.2 Label Propagation Algorithm (标签传播算法)

1. Compute affinity matrix W (亲和矩阵).
2. Compute degree matrix D , where $D_{ii} = \sum_j W_{ij}$.
3. Initialize $\hat{Y}^{(0)} \leftarrow (y_1, \dots, y_l, 0, 0, \dots, 0)$.
4. Iterate: $\hat{Y}^{(t+1)} = D^{-1}W\hat{Y}^{(t)}$, fix labels of labeled data.
5. Repeat until convergence.
6. Assign labels based on the sign of $\hat{y}_i^{(\infty)}$.

- The algorithm forces the labels on the labeled data (强制在标记数据上保留标签)
- The algorithm tries to maximizes the consistency of the unlabeled examples with the topology of the graph (试图使未标记示例与图的拓扑结构保持最大程度的一致性)

其实到最后也不知道这一讲到底细讲了什么，并没有什么公式(雾)

6 Lec5: GNN

to be continued...