

# Note and Formula of DDA4210

softstar

January 27, 2026

## 1 Lec1: introduction

There are no formulas in Lecture 1. Good luck with the rest!

## 2 Lec2: Advanced Ensemble Learning

### 2.1 Gradient Boosting:

Boosting is an ensemble technique that combines multiple weak learners to create a strong learner. The main idea is to train models sequentially, with each model focusing on the errors made by the previous ones. (用人话说，就是一波接一波地训练模型，每一波都专注于纠正前一波的错误，从而逐步提升整体的预测能力。)

- **Final Model:**

$$H_T(\mathbf{x}) = \sum_{t=1}^T \alpha_t h_t(\mathbf{x})$$

- The weighted sum of all weak learners(加权和).
- $H_T(\mathbf{x})$  is the final model after  $T$  rounds.
- $h_t(\mathbf{x})$  is the  $t$ -th weak learner.
- $\alpha_t$  is the weight of the  $t$ -th weak learner.

- **Loss Function:**

$$\mathcal{L}(H) := \frac{1}{n} \sum_{i=1}^n \ell(H(\mathbf{x}_i), y_i)$$

- **Each Step:** We want to add a function  $h$  to minimize the loss as fast as possible. Using first-order Taylor expansion(一阶泰勒展开):

$$\mathcal{L}(H + \alpha h) \approx \mathcal{L}(H) + \alpha \langle \nabla \mathcal{L}(H), h \rangle$$

- **Find  $h$ :** Minimize  $\langle \nabla \mathcal{L}(H), h \rangle$ , i.e.:

$$h = \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n \frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]} h(\mathbf{x}_i)$$

\* Here  $\frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]}$  is the gradient of the loss function w.r.t. the current model's output on the  $i$ -th sample.

\* We train  $h$  to fit these **negative gradients**(负梯度).

- **GBR with squared error loss:** For regression tasks with squared error loss:

$$\ell(H(\mathbf{x}_i), y_i) = \sum_{i=1}^n (y_i - H(\mathbf{x}_i))^2$$

- Solve  $h_{t+1} = \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n q_i h(\mathbf{x}_i)$ , where  $q_i = \frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]}$
- Let  $\sum_{i=1}^n h^2(\mathbf{x}_i) = \text{constant}$  (we can always normalize the predictions(对结果归一化处理)) and replace  $q_i$  with  $-2r_i$ . We have

$$\begin{aligned} h_{t+1} &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n q_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} -2 \sum_{i=1}^n r_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n (r_i^2 - 2r_i h(\mathbf{x}_i) + (h(\mathbf{x}_i))^2) \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n (h(\mathbf{x}_i) - r_i)^2 \end{aligned} \tag{1}$$

- We train  $h_{t+1}$  to predict  $r_i$ , which are from the old model  $H_t$ .

- The gradient is:

$$\frac{\partial \mathcal{L}}{\partial [H(\mathbf{x}_i)]} = -2(y_i - H(\mathbf{x}_i))$$

- So we fit  $h$  to the residuals:

$$r_i = y_i - H(\mathbf{x}_i)$$

- Update the model:

$$H_t(\mathbf{x}) = H_{t-1}(\mathbf{x}) + \alpha_t h_t(\mathbf{x})$$

- **GBR with Absolute loss (更鲁棒):** For regression tasks with absolute loss:

- Square loss is easy to deal with mathematically but not robust to outliers.
- Absolute loss (more robust to outliers):

$$\ell(y, \hat{y}) = |y - \hat{y}|$$

- The gradient is  $q_i = \frac{\partial \mathcal{L}}{\partial H(\mathbf{x}_i)} = -\text{sign}(y_i - H(\mathbf{x}_i))$ .
- Then fit  $h$  on  $-q_i$ ,  $i = 1, 2, \dots, n$ . (no longer the residuals, different from using the squared loss)

- **GBR with Huber loss(更更鲁棒):**

- Huber loss (more robust to outliers):

$$\ell(y, \hat{y}) = \begin{cases} \frac{1}{2}(y - \hat{y})^2 & |y - \hat{y}| \leq \delta \\ \delta(|y - \hat{y}| - \delta/2) & |y - \hat{y}| > \delta \end{cases}$$

- The gradient is:

$$\frac{\partial \mathcal{L}}{\partial H(\mathbf{x}_i)} = \begin{cases} -(y_i - H(\mathbf{x}_i)) & |y_i - H(\mathbf{x}_i)| \leq \delta \\ -\delta \operatorname{sign}(y_i - H(\mathbf{x}_i)) & |y_i - H(\mathbf{x}_i)| > \delta \end{cases}$$

- **GBM for classification:**

- Predict probability of  $K$  classes:

$$p_k(\mathbf{x}) = \frac{\exp(h^{(k)}(\mathbf{x}))}{\sum_{c=1}^K \exp(h^{(c)}(\mathbf{x}))} \triangleq \hat{y}^{(k)}, \quad k = 1, 2, \dots, K$$

- Loss:  $\mathcal{L}(H) = \sum_{i=1}^n \ell(\mathbf{y}_i, \hat{\mathbf{y}}_i)$  (e.g. cross-entropy or KL divergence)

- Initialize  $H^{(1)}, H^{(2)}, \dots, H^{(K)}$ , iterate until converge or reach max  $T$ :

1. Calculate negative gradients for every class:

$$-g_k(\mathbf{x}_i) = -\frac{\partial \mathcal{L}}{\partial [H^{(k)}(\mathbf{x}_i)]}, \quad i = 1, \dots, n, \quad k = 1, \dots, K$$

2. Fit  $h^{(k)}$  to  $-g_k(\mathbf{x}_i)$  (负梯度),  $k = 1, 2, \dots, K$ .

3. Update:  $H^{(k)} \leftarrow H^{(k)} + \alpha h^{(k)}$ ,  $k = 1, 2, \dots, K$ .

## 2.2 AdaBoost

A special case of gradient boosting with exponential loss.

- Exponential loss (learns  $\alpha$  adaptively):

$$\mathcal{L}(H) = \sum_{i=1}^n e^{-y_i H(\mathbf{x}_i)}$$

- Gradient:  $q_i = \frac{\partial \mathcal{L}}{\partial H(\mathbf{x}_i)} = -y_i e^{-y_i H(\mathbf{x}_i)}$
- Let  $w_i = \frac{1}{Z} e^{-y_i H(\mathbf{x}_i)}$ , where  $Z = \sum_{i=1}^n e^{-y_i H(\mathbf{x}_i)}$  (constant w.r.t  $h$ ), so  $\sum_{i=1}^n w_i = 1$ .  $w_i$  is the relative contribution of  $(\mathbf{x}_i, y_i)$  to the overall loss.
- Binary classification:  $y \in \{-1, +1\}$ ,  $h(\mathbf{x}) \in \{-1, +1\}$ .
- Derivation:

$$\begin{aligned} h &= \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n q_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} - \sum_{i=1}^n w_i y_i h(\mathbf{x}_i) \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i - \sum_{i:h(\mathbf{x}_i) = y_i} w_i \\ &= \arg \min_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i \end{aligned}$$

\* Last equality holds because  $\sum_{i=1}^n w_i = 1$ .

- Result:

$$h = \arg \min_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i$$

$$\epsilon := \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i$$

is the weighted classification error(加权错误率).

Note: misclassified points by  $H$  get larger weights(分类错误的点会得到更大的权重).

- Given  $h$ , find  $\alpha$  via:

$$\alpha = \arg \min_{\alpha} \mathcal{L}(H + \alpha h) = \arg \min_{\alpha} \sum_{i=1}^n e^{-y_i(H(\mathbf{x}_i) + \alpha h(\mathbf{x}_i))}$$

- Differentiate w.r.t  $\alpha$  and set to zero:

$$\sum_{i=1}^n y_i h(\mathbf{x}_i) e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} = 0$$

It follows that:

$$\begin{aligned} \sum_{i:h(\mathbf{x}_i)y_i=1} e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} - \sum_{i:h(\mathbf{x}_i)y_i=-1} e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} &= 0 \\ \sum_{i:h(\mathbf{x}_i)y_i=1} w_i e^{-\alpha} - \sum_{i:h(\mathbf{x}_i)y_i=-1} w_i e^{\alpha} &= 0 \end{aligned}$$

We have  $(1 - \epsilon)e^{-\alpha} - \epsilon e^{\alpha} = 0$ ,  $e^{2\alpha} = \frac{1-\epsilon}{\epsilon}$ , and get:

$$\boxed{\alpha = \frac{1}{2} \ln \frac{1-\epsilon}{\epsilon}}$$

## 2.3 Mixture of Experts(MoE)

A machine learning technique where multiple expert learners (e.g. neural networks) are used to divide a problem space into homogeneous regions (distinct subtasks). (用人话 (AI) 说就是，把一个复杂的问题拆分成多个子任务，每个子任务由一个专家模型来处理，从而提升整体的学习效果。)

(骗你的人话也没看懂。。)

### 2.3.1 The First Attempt

- Error function:

$$E = \left\| y - \sum_{j=1}^k g_j O_j \right\|^2$$

- $y$ : target vector;  $O_j$ : output of expert  $j$ ;  $g_j$ : proportional contribution of expert  $j$ .

- \*This error function does not ensure localisation of experts.(人话：专家没有明确的分工).

### 2.3.2 The Second Attempt

- Error function:

$$E = \sum_{j=1}^k g_j \|y - O_j\|^2$$

- The system tends to devote a single expert to each training case.
- \*This may not work well in practice.(人话: 实际效果可能不佳)

### 2.3.3 The Third Attempt [Jacobs et al. 1991]

- Error function (mixture of Gaussians):

$$E_{ME} = -\log \sum_{j=1}^k g_j \exp \left( -\frac{1}{2}(y - O_j)^T \Sigma^{-1}(y - O_j) \right)$$

- Assume  $\Sigma = I$  ( $\Sigma$ : Covariance matrix, 协方差矩阵), derivative w.r.t the  $i$ -th expert:

$$\frac{\partial E_{ME}}{\partial O_i} = - \left[ \frac{g_i \exp(-\frac{1}{2}(y - O_i)^T(y - O_i))}{\sum_j g_j \exp(-\frac{1}{2}(y - O_j)^T(y - O_j))} \right] (y - O_i)$$

- Compare with derivative of second attempt:

$$\frac{\partial E}{\partial O_i} = -2g_i(y - O_i)$$

- The former considers how well expert  $i$  performs relative to others, adapting the best-fitting expert faster.

E.g.,  $g_1 = 0.8$ ,  $g_2 = 0.2$ , then  $\frac{0.8 \times 0.9}{0.8 \times 0.9 + 0.2 \times 0.1} \approx 0.97 > 0.8$ .  
(人话: 表现好的专家会得到更快的提升)

(这 nm 都是什么玩意?)

## 2.4 Stacking(堆叠法)

Multiple base learners' outputs  $(\hat{y}_1, \hat{y}_2, \dots, \hat{y}_N)$  are combined by a meta-learner (stacker) to produce the final prediction  $\hat{Y}$ . (用人话说就是, 把多个模型的预测结果再拿去训练一个新的模型, 从而提升整体的预测效果。)

- Multi-level stacking: stacking can be repeated.
- Popular stackers: linear models (fast), gradient boosting (accurate).
- Base models should be diverse, expert at different data parts.
- Trade-off: accurate but slow to predict.

(为什么一讲能有这么多神秘小知识点和神秘公式? )

### 3 Lec3: Advanced Applications

#### 3.1 Recommendation Systems

##### 3.1.1 Collaborative Filtering(协同过滤)

Core idea: Use behavioral data from many users (e.g., ratings, clicks) to predict what the current user might like. (人话: 大数据推荐你喜欢的东西)

- **User-Item Interaction:**
  - Explicit Feedback(显式反馈): ratings, purchases.
  - Implicit Feedback(隐式反馈): clicks, browsing time.
- **User-Item Rating Matrix:** Rows = users, columns = items, values = ratings. Typically large and sparse.
- **Distance/Similarity Measurement(相似度度量):**
  - Euclidean distance:  $\text{sim}(\text{user}_i, \text{user}_j) = \frac{1}{1 + \|\mathbf{x}_i - \mathbf{x}_j\|_2}$
  - Cosine similarity:  $\text{sim}(\text{user}_i, \text{user}_j) = \frac{\mathbf{x}_i \cdot \mathbf{x}_j}{\|\mathbf{x}_i\| \|\mathbf{x}_j\|}$
  - Pearson correlation:  $\rho_{X,Y} = \frac{\text{cov}(X,Y)}{\sigma_X \sigma_Y}$
- **Nearest-Neighbor Collaborative Filtering(最近邻协同过滤):**  
Predict utility of item  $i$  based on similar users who rated that item.
  - $\mathcal{N}$ : neighborhood set (most similar users to  $u$  who rated item  $i$ ) (人话: 给你推荐东西的那些“邻居”用户)
  - $w_{uv} \in [0, 1]$ : similarity weight between users  $u$  and  $v$  (人话: 你和邻居用户的相似度)
  - Prediction:  

$$\hat{x}_{ui} = \bar{x}_u + \sum_{v \in \mathcal{N}} \left( (x_{vi} - \bar{x}_v) \times \frac{w_{uv}}{\sum_{v' \in \mathcal{N}} w_{uv'}} \right)$$
    - $\bar{x}_u = \frac{1}{|I_u|} \sum_{i \in I_u} x_{ui}$  (average rating of user  $u$ , where  $I_u$  is the set of items rated by  $u$ )
    - $\bar{x}_v = \frac{1}{|I_v|} \sum_{i \in I_v} x_{vi}$  (average rating of user  $v$ , where  $I_v$  is the set of items rated by  $v$ )
- **Matrix Factorization Collaborative Filtering(矩阵分解协同过滤):**

- Notations:

- \*  $R = [r_{ui}] \in \mathbb{R}^{m \times n}$ : incomplete user-item rating matrix
- \*  $\Omega$ : set of observed entries (known ratings)
- \*  $P = [p_1, \dots, p_u, \dots, p_m] \in \mathbb{R}^{f \times m}$ ,     $Q = [q_1, \dots, q_i, \dots, q_n] \in \mathbb{R}^{f \times n}$

- Basic SVD ( $R \approx P^T Q$ ):

$$\min_{P,Q} \sum_{(u,i) \in \Omega} \{(r_{ui} - p_u^T q_i)^2 + \lambda(\|p_u\|^2 + \|q_i\|^2)\}$$

- SVD with bias ( $b_{ui} = \mu + b_u + b_i$ ):

$$\min_{P,Q,B} \sum_{(u,i) \in \Omega} \{(r_{ui} - \mu - b_u - b_i - p_u^T q_i)^2 + \lambda(\|p_u\|^2 + \|q_i\|^2 + b_u^2 + b_i^2)\}$$

- $\mu$ : global mean;  $b_u$ : user bias;  $b_i$ : item bias;
- $\lambda(\cdot)$ : regularization term (prevents overfitting).
- Optimization: GD/SGD or Alternating Least Squares(交替最小二乘法).

- **Pros & Cons of Collaborative Filtering:**

- Pros: No domain knowledge needed; captures diverse user preferences(人话: 不需要领域知识, 能捕捉多样的用户偏好).
- Cons: Cold-start problem (new users/items); data sparsity(人话: 新来的用户和物品数据较少, 导致推荐效果差).

### 3.1.2 Content-Based Methods

- **Content analysis:** item  $\rightarrow$  feature vector  $v$  (e.g. TF-IDF, image features).
- **Profile learning:** user  $\rightarrow$  feature vector  $z$  (e.g. age, sex, education).
- **Filtering module:** train classification/regression model to predict user's utility for an item.
- **Recommendation for user:**
  - $n$ : # of items;
  - $z_u \in \mathbb{R}^d$ :
  - user  $u$ 's feature vector;
  - $h : \mathbb{R}^d \rightarrow \mathbb{R}^n$  (e.g. neural network);
  - $h_i$ :  $i$ -th output of  $h$ .
  - $\ell$ : loss function (e.g. squared loss).

$$\min_h \sum_{(u,i) \in \Omega} \ell(r_{ui}, h_i(z_u))$$

- **Recommendation for item:**

- $m$ : # of users;
- $v_i \in \mathbb{R}^{d'}$ : item
- $i$ 's feature vector;
- $g : \mathbb{R}^{d'} \rightarrow \mathbb{R}^m$  (e.g. neural network);
- $g_u$ :  $u$ -th output of  $g$ .
- $\ell$ : loss function (e.g. squared loss).

$$\min_g \sum_{(u,i) \in \Omega} \ell(r_{ui}, g_u(v_i))$$

- **Pros & Cons of Content-Based Methods:**

- Pros: User-independent; explainable; handles new items/users well.
- Cons: Needs domain knowledge; narrow recommendations (similar items).

### 3.1.3 Hybrid Methods(看起来是不重要的知识点)

Most modern systems are hybrid recommenders.

- Combine separate recommenders (CF + CB): ensemble techniques (linear weighting, stacking, etc.)
- Add content-based aspects to CF: e.g. matrix factorization with side information.

### 3.1.4 Evaluation Metrics for RS

#### 3.1.4.1 Prediction Metrics(评分预测指标)

- **Mean Absolute Error (MAE):**

$$\text{MAE} = \frac{1}{|\mathcal{T}|} \sum_{(u,i) \in \mathcal{T}} |r_{ui} - \hat{r}_{ui}|$$

- **Root Mean Squared Error (RMSE):**

$$\text{RMSE} = \sqrt{\frac{1}{|\mathcal{T}|} \sum_{(u,i) \in \mathcal{T}} (r_{ui} - \hat{r}_{ui})^2}$$

Here  $\mathcal{T}$  denotes the set of user–item pairs used for evaluation (e.g., the test set).(人话：测试集中所有用户和物品的组合)

#### 3.1.4.2 Ranking-based Metrics(基于排序的指标)

- **Precision@K:** fraction of top- $K$  recommended items that are relevant(人话：前 K 个推荐中有多大是相关的).

$$\text{Prec}(R)_k = \frac{|\{r \in R : r \leq k\}|}{k}$$

- **Recall@K:** fraction of relevant items covered in top- $K$ (人话：前 K 个推荐覆盖了多少相关的物品).

$$\text{Recall}(R)_k = \frac{|\{r \in R : r \leq k\}|}{|R|}$$

- (Precision = TP/(TP+FP); Recall = TP/(TP+FN)).
- $r$  = rank position of a recommended item;
- $k$  = cut-off (top- $k$ );
- $R$  = set of relevant items for the user.

- **Average Precision (AP@N):** average of precision values at ranks of relevant items.(前 N 个推荐中相关物品的精确率)

$$\text{AP}@N = \frac{1}{m} \sum_{k=1}^N P(k) \cdot \text{rel}(k)$$

where  $P(k)$  is precision@ $k$ ,  $m$  number of relevant items, and  $\text{rel}(k)$  is indicator if item at rank  $k$  is relevant.

- **Mean Average Precision (MAP):** mean of AP over  $Q$  users:

$$\text{MAP} = \frac{1}{Q} \sum_{q=1}^Q \text{AP}(q)$$

- **Normalized Discounted Cumulative Gain (NDCG):** evaluates ranked relevance with position discounting.

$$NDCG_p = \frac{DCG_p}{IDCG_p}, \quad DCG_p = \sum_{i=1}^p \frac{2^{rel_i} - 1}{\log_2(i + 1)}$$

where  $rel_i$  is relevance of item at rank  $i$ , and  $IDCG_p$  is the ideal DCG (sorted by relevance).

Range:  $[0, 1]$ . (NDCG 越接近 1 越好)

- **Example:** 5 recommended items with relevances  $[3, 2, 1, 0, 2]$  (in rank order).

$$DCG_5 = \frac{2^3 - 1}{\log_2(1 + 1)} + \frac{2^2 - 1}{\log_2(2 + 1)} + \frac{2^1 - 1}{\log_2(3 + 1)} + \frac{2^0 - 1}{\log_2(4 + 1)} + \frac{2^2 - 1}{\log_2(5 + 1)} \approx 10.5538,$$

$$IDCG_5 = DCG_5(\text{sorted rel} = [3, 2, 2, 1, 0]) \approx 10.8235,$$

$$NDCG_5 = \frac{DCG_5}{IDCG_5} \approx 0.975.$$

(这一坨又是什么玩意？)

### 3.2 Learning to Rank(L2R) (排序学习)

Learning to Rank (L2R) trains models to order items by relevance, optimizing ranking-specific objectives such as pairwise or listwise losses. (你说得对，但是这里好像也没有什么公式啊(雾))

- L2R is a supervised learning problem for ranking.
- Training data consists of:
  - A set of queries  $Q = \{q_1, \dots, q_m\}$
  - A set of documents  $D$
  - For each query  $i$ , relevant documents  $D_i = \{d_{i,1}, \dots, d_{i,n_i}\} \subseteq D$
  - Relevance scores  $\mathbf{y}_i = (y_{i,1}, \dots, y_{i,n_i})$  for each  $d_{i,j}$
- Goal: Given a new query  $q$ , output a sorted list of relevant documents.
- **Point-wise Modeling:** Predicts each (query, document) pair independently.
 

*Pro:* Simple, can use regression/classification.  
*Con:* Ignores relative order between documents.  
 (一句话：点对点建模简单但忽略了文档间的相对顺序。)
- **Pair-wise Modeling:** Predicts preference between document pairs for a query.
 

*Pro:* Models relative order.  
*Con:* Cannot distinguish excellent-bad from fair-bad pairs.  
 (一句话：成对建模能捕捉相对顺序，但无法区分优秀-差和一般-差的对比。)
- **List-wise Modeling:** Predicts for the whole ranked list of documents.
 

*Pro:* Considers position in ranking, aligns with ranking metrics.  
*Con:* High training complexity.  
 (一句话：列表建模考虑排名位置，但训练复杂度高。)
- **Evaluation for L2R:** Use benchmark datasets and ranking metrics (e.g., MAP, NDCG).

- Algorithms for L2R:

- Example: Ranking SVM (pairwise):

- \* **Goal:** Learn a scoring function  $h(x) = w^\top x$  such that for any pair with labels  $y_i > y_j$  we have  $h(x_i) > h(x_j)$ . (人话: 让相关性更高的文档得分更高)

- \* **Training pairs:** Construct pair set  $\mathcal{P} = \{(i, j) : y_i > y_j\}$ ;  $m = |\mathcal{P}|$  denotes number of pairs.

- \* **Optimization (primal):**

$$\begin{aligned} \min_{w, \xi_{ij} \geq 0} \quad & \frac{1}{2} \|w\|^2 + \frac{C}{m} \sum_{(i,j) \in \mathcal{P}} \xi_{ij} \\ \text{s.t.} \quad & w^\top x_i \geq w^\top x_j + 1 - \xi_{ij}, \quad \forall (i, j) \in \mathcal{P}. \end{aligned}$$

- \* **Notes:**  $\xi_{ij}$  are hinge-loss slacks;  $C$  controls margin vs. training error. The objective is equivalent to minimizing the average pairwise hinge loss.

- \* **Prediction:** Score each document by  $h(x) = w^\top x$  and sort descending to produce a ranking.

- \* **Remarks:** Works well for pairwise preferences; training can be expensive due to  $O(n^2)$  pairs, so sampling or stochastic methods are often used. Kernel SVMs and regularization extend naturally.

- \* (人话总结: Ranking SVM 通过学习一个线性评分函数来排序文档, 优化目标是最大化正确排序对的间隔 + 最小化排序错误的惩罚。训练时需要处理大量文档对, 复杂度高。)

(叽里咕噜说什么在)

## 4 Lec4-1: Graph Cut and Spectral Clustering(谱聚类)

### 4.1 Graph Partition

A similarity graph  $G=(V,E,W)$  represents data points as vertices  $V$ , with an edge in  $E$  when the pairwise similarity is positive and weights  $W$  storing those affinities. The affinity matrix records these pairwise similarities. Graph partitioning (clustering) aims to split the graph so that edges inside a group have large weights while edges across groups have small weights. (人话: 图划分就是把图分成若干部分, 使得每个部分内的节点之间联系紧密 (组内权重大), 而不同部分之间的联系较弱 (组间权重小)。)

Given data points, a similarity graph can be constructed using methods such as **k-nearest neighbor** or  $\epsilon$ -neighborhood. The edge weights are often defined by a **Gaussian kernel**:

$$k(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

### 4.2 Minimum Cut

- Minimum cut: partition the graph into two sets  $A$  and  $B$  minimizing  $\text{cut}(A, B) := \sum_{i \in A, j \in B} w_{ij}$ .
- Solvable efficiently (e.g. via max-flow/min-cut algorithms, typical cost  $O(|V||E|)$ ), but the minimum cut often yields unbalanced solutions (it may isolate vertices). (人话: 最小割可以高效求解, 但结果往往不平衡, 可能会把一些节点孤立出来。)

- Not satisfactory partition? Often isolates vertices

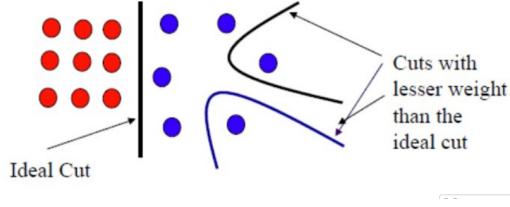


图 1: Minimum Cut Example

(画的不错的一张图)

- To address this, we can use **Normalized Cut (Ncut)**:

### 4.3 Normalized Cut

Normalized Cut balances cut weight with cluster sizes. For a partition  $(A, B)$ , define

$$\text{vol}(A) = \sum_{i \in A} d_i, \quad d_i = \sum_j w_{ij},$$

and

$$\text{Ncut}(A, B) := \text{cut}(A, B) \left( \frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right).$$

- Minimizing Ncut favors balanced partitions but is **NP-hard**;
- **spectral clustering (谱聚类)** provides an efficient relaxation.

#### 4.3.1 Degree Matrix and Graph Laplacian

$$W = \begin{bmatrix} w_{11} & w_{12} & \cdots & w_{1N} \\ w_{21} & w_{22} & \cdots & w_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ w_{N1} & w_{N2} & \cdots & w_{NN} \end{bmatrix}$$

$$D = \begin{bmatrix} d_1 & 0 & \cdots & 0 \\ 0 & d_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & d_N \end{bmatrix}, \quad d_j = \sum_{i=1}^N w_{ij}.$$

- $D$  is the **degree matrix**. And the **graph Laplacian matrix** is defined as

$$L := D - W.$$

- Properties:  $L$  is symmetric positive semi-definite,  $\mathbf{1}^\top L = 0$ , and its eigenvectors are used in spectral clustering (relaxation of Ncut).

### 4.3.2 Normalized Cut and Graph Laplacian

#### Mathematical derivation (optional)

Recall  $L = D - W$  and  $D = \text{diag}(d_1, \dots, d_N)$ .

Let  $u = [u_1, u_2, \dots, u_N]^\top$  with

$$u_i = \begin{cases} \frac{1}{\text{vol}(A)}, & \text{if } i \in A, \\ -\frac{1}{\text{vol}(B)}, & \text{if } i \in B. \end{cases}$$

Then

$$u^\top Lu = \frac{1}{2} \sum_{i,j} w_{ij} (u_i - u_j)^2 = \sum_{i \in A, j \in B} w_{ij} \left( \frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right)$$

and

$$u^\top Du = \sum_i d_i u_i^2 = \sum_{i \in A} \frac{d_i}{\text{vol}(A)^2} + \sum_{j \in B} \frac{d_j}{\text{vol}(B)^2} = \frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)}.$$

Therefore

$$\frac{u^\top Lu}{u^\top Du} = \sum_{i \in A, j \in B} w_{ij} \left( \frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right) = \text{Ncut}(A, B).$$

(这几把都是什么玩意？好像是上面的推导吧)

(不管了，反正是 optional，摆在这图个吉利)

#### Conclusion

- Ncut is equivalent to minimizing the Rayleigh quotient  $\frac{u^\top Lu}{u^\top Du}$ , i.e.,

$$\min_{A,B} \text{Ncut}(A, B) \iff \min_u \frac{u^\top Lu}{u^\top Du}, \quad u \in \mathbb{R}^N, \quad u_i = \begin{cases} \frac{1}{\text{vol}(A)}, & i \in A, \\ -\frac{1}{\text{vol}(B)}, & i \in B. \end{cases}$$

- Equivalent formulation: minimize the quotient subject to  $u^\top D\mathbf{1} = 0$  and binary constraints  $u_i \in \{1, -b\}$  (for some  $b > 0$ ).
- Relaxation:

$$Lu = \lambda Du,$$

taking the eigenvector corresponding to the **second smallest eigenvalue** as the relaxed solution.

- Equivalently, use the normalized Laplacian  $\tilde{L} = D^{-1}L = I - D^{-1}W$ .  
Obtain a binary partition by thresholding  $u$  at 0:  $i \in A$  if  $u_i \geq 0$ , else  $i \in B$ .
- Extend to  $k$  clusters by using the first  $k$  nontrivial eigenvectors and applying  $k$ -means (spectral clustering).
- (人话：2类划分找第二小特征值对应的特征向量， $k$ 类划分用前  $k$  个非平凡特征向量，然后  $k$ -means 聚类。)

(事实上还是没看懂，插个眼以后复习的时候看看有没有什么需要补的)

#### 4.4 Spectral Clustering Algorithm

- Input: data  $X = \{x_1, x_2, \dots, x_N\}$  and number  $K$  of clusters.
- **Step 1:** Construct a similarity (affinity) matrix  $W$  (e.g. Gaussian kernel)

$$w_{ij} = \exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma^2}\right)$$

and build either a  $k$ -nearest neighbor or  $\epsilon$ -neighborhood graph.

- **Step 2:** Compute Laplacian  $L$  (or normalized variants):

$$L = D - W, \quad \tilde{L} = D^{-1}L = I - D^{-1}W, \quad \hat{L}_{sym} = I - D^{-1/2}WD^{-1/2}.$$

- **Step 3:** Eigen-decompose (normalized) Laplacian and take first  $K$  nontrivial eigenvectors to form  $Z = [v_1, \dots, v_K]^\top \in \mathbb{R}^{K \times N}$ .(对 L 做特征值分解)
- **Step 4:** Normalize the columns (row-wise embedding) to unit  $\ell_2$  norm(归一化):

$$z_i \leftarrow z_i / \|z_i\|, \quad i = 1, \dots, N.$$

- **Step 5:** Run  $K$ -means on  $\{z_1, \dots, z_N\}$  and output  $K$  clusters (assignments on  $Z$  or map back to  $X$ ).
- Notes: use  $\tilde{L}_{sym}$  (symmetric normalized) for best numerical stability; thresholding the second eigenvector recovers a binary partition.

- **Properties of L:**

For  $L = D - W$  or  $\hat{L} = I - D^{-1/2}WD^{-1/2}$

- $L$  (and  $\hat{L}$ ) are symmetric and positive semi-definite.(对称 + 半正定)
- Eigenvalues satisfy  $0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_N$ .
- The multiplicity  $K$  of eigenvalue 0 equals the number of connected components of the graph (hence  $K$  clusters). (0 特征值的数量等于图的连通分量数, 也就是聚类数)

## 5 Lec4-2: Semi-Supervised Learning

to be continued...